

Zepto users already reach 18 categories. The problem is the month they stop.

Getting more monthly actives into a new category — by fixing the thing Zepto's own speed broke.

47.97M

annual transacting users, FY26


₹59.40

lost per order, Q4FY26

29% → 44%

non-grocery share of q-commerce
GMV, CY25→CY30

Breadth isn't the problem. The rate of new categories per month is — and it has already fallen 30%.



Months 1 – 12

0.78

new categories / month

Months 12 – 23

0.55


a 30% collapse in rate

Why the obvious answer is wrong

- “Users are stuck in a habit loop, so show them more.” But they add 15 categories in two years. They are not stuck.
- The goal is a monthly rate, not a lifetime total. A flattening cumulative curve is a rate approaching zero.
- So the question is what is different about the categories still left after month 12.


Three points are disclosed in Zepto's DRHP (2.6 at month 1, 12+ at month 12, 18.0 at month 23). The rates are arithmetic on those figures; the intermediate curve shape is not disclosed and is not assumed.

A six-stage pipeline that reads 3,000 reviews and cites a checkable source for every claim.



A live, runnable version of every stage — including a classifier you can paste any review into — is linked on the final slide.

3,000 reviews cannot tell you why someone didn't buy. Establishing that is the finding.



Zero fit mentions. Four competitor mentions in 628 reviews, none of them Nykaa. Absence at this scale is evidence about the source, not about users.

My first run said trust was 71% of all barriers. It was measuring the word “review”.

The bug my own audit caught

Run 1: 71% trust

A clean, confident, completely wrong chart. Nothing about it looked broken.

Provenance check

I made the classifier report which term fired on every row. “Review” and “rating” were matching self-reference — “writing this review”, “giving 1 star rating” — which says nothing about trusting a product.

Two real bugs

Trust cues now need product context. And all matching moved to word boundaries, because “cat” was matching inside “category”, making it the top category.

Run 2: 38%, and different

Trust dropped and its meaning changed: expiry, not reviews. Assortment overtook it.

An unaudited classifier looks identical whether or not it is right.

Scored against a human

30 reviews, stratified by label, fixed seed, hand-labelled blind before seeing any model output.

0.32

Cohen's kappa — fair, not substantial.
Reported as-is.

53%

raw agreement across 7 classes

0/20

negative-control leakage — the gate isn't inventing signal

The errors run one way

10 of 30 disagreements are the model finding a barrier where I saw none; only 3 are the reverse. It over-triggers, so 107 is an upper bound and the real signal is thinner still.

Zepto is a retrieval interface being asked to sell deliberation. The public data can't see it — which is why I asked users.

ESTABLISHED

The rate collapses

Zepto's own cohort disclosure: 0.78 new categories per month in year one, 0.55 in year two. Arithmetic, not inference. The problem is real and it is getting worse.

RULED OUT

It isn't awareness, and it isn't visible in reviews

3 awareness mentions in 3,000 reviews. What surfaces instead is stock-outs and expiry, both operational. Public review data is the wrong instrument for a question about purchases that never happened.

FOUND IN INTERVIEWS

The product page cannot answer the question

3 of 6 named the same gap unprompted: no verified reviews, stock images, no specs, no seller verification. Two go to Nykaa and Amazon for exactly that. A fourth never recalls that beauty exists on the app.

THE SEGMENT

The Plateaued Regular

- 6+ orders a month
- 12+ months on the app
- A fixed 6–9 category repertoire
- No new category in 60+ days

THE ROOT CAUSE

The ten-minute promise sets the user's mental mode to task execution, and the app has no mechanism for resolving uncertainty. The categories left unentered after month 12 are the ones carrying personal consequence — sunscreen, a first protein powder, a new puppy's food.

Zepto's greatest strength is the direct cause of the metric it is now trying to move.

This is not a growth metric. It is the only margin lever left that does not require buying more users.

49.54M → 47.97M

Annual transacting users fell sequentially in Q4FY26 while quarterly orders rose from 167M to 210M. Acquisition is slowing; growth now has to come from existing users going wider.

₹59.40


Lost per order in Q4FY26. Blinkit made ₹1.35 on theirs. Every incremental order in a low-margin staple deepens the hole.

29% → 39–44%

Non-grocery share of quick-commerce GMV, CY25 to CY30. The higher-margin categories — and the ones a retrieval interface cannot sell.

So the categories Zepto needs for margin are the categories its interface is worst at selling. Fixing deliberation unlocks both problems at once.

First Buy: one observation, one question, one answer — in the cart, because that is where users said yes.



SURFACE

The cart — and the interviews are why it moved there

I designed this for the post-payment tracking screen. No user gave me a reason to believe that works. Two had already converted via a checkout prompt, one asked for cart bundling by name, and nobody expanded by browsing.

TRIGGER

An observation about their own behaviour, not a recommendation

“You've ordered cereal and milk 21 times, never fruit.” Specific and checkable, so it reads as the app noticing rather than the app advertising.

RESOLUTION

One question, then one product — carrying what the PDP is missing

Users named the gap precisely: verified reviews, real photos, dimensions, seller verification, ingredients. So the card carries a rating summary, one reason from mined review language, and the top objection pre-handled.

DE-RISK

Free return on the first buy in any new category

Attacks the risk structurally rather than with a discount. Cheaper per conversion, and it does not train users to wait for coupons.

Ranking is not the bottleneck — collaborative filtering solved that years ago. Justification is. That is the generative problem, and it is fed by the corpus that diagnosed the barrier.

How I would know it worked — including the one metric that could prove it did not.

North star	% of MAC buying from ≥ 1 new category per month	The stated goal. Needs two monthly cycles, so it reads as directional at four weeks.
Primary	New-category purchases per 1,000 cart-card impressions	What the experiment is actually powered on.
Secondary	30-day repeat rate within the newly entered category	Separates behaviour change from a novelty transaction. Buy sunscreen once and never again and this made a sale, not a habit.
Guardrail	Return rate on First Buy items	Free returns must not quietly become a loss centre.
Guardrail	Cart-to-order conversion	The hard one. A card in the cart must never cost an order that was already converting.
Guardrail	Card dismiss rate	An annoyance signal, and the trigger for the frequency cap.

Experiment: geo-split A/B on the Plateaued Regular cohort, four weeks, powered on the primary metric. Stopping rule set in advance — if cart-to-order conversion drops beyond the pre-agreed margin, the test ends early regardless of how the primary metric is doing.
